

LOAD TESTING YOUR APP

CONFOO MONTREAL 2020

Ian Littman / @iansltx

follow along at <https://ian.im/loadfoo20>

SPEED.

SCALABILITY.

STABILITY.

QUESTIONS WE'LL ANSWER

- What types of tests exist, and what sets each type apart?
- When should I build and run performance tests?
- **How can I match my load test with (anticipated) reality?**
- What does a real load test script look like on a small system?
- How do I properly analyze results during and after my test?

QUESTIONS WE WON'T ANSWER

- How do I use \$otherPerfTestTool (!== 'k6')?
- How can I set up clustered load testing?
- How can I simulate far-end users?
- How can I test web page performance browser-side?
- How can I do deep application profiling? (Blackfire for PHP)
- What about single-user load testing?

WE'LL BE TESTING WITH K6*

- Write your tests in JS**
- Run via a Go binary***
- HAR import for in-browser recording

* More tools are listed at the end of this presentation.

** Uses goja, not V8 or Node, and doesn't have a global event loop yet.

*** I've used this on a project significantly more real than the example in this presentation, so that's a big reason we're looking at it today.

#IFDEF

- Smoke Test
- Load Test vs. Stress Test
- Soak Test vs. Spike Test

SMOKE TEST

- An initial test to confirm the system operates properly without a large amount of generated load
- Do this before you load test
- Pick your implementation...
 - Integration tests in your existing test suite
 - Load test script, turned down to one (thorough) iteration and one **Virtual User** (VU)

LOAD TEST

- \leq expected peak traffic
- Your system shouldn't break
- If it does, it's a...

STRESS TEST

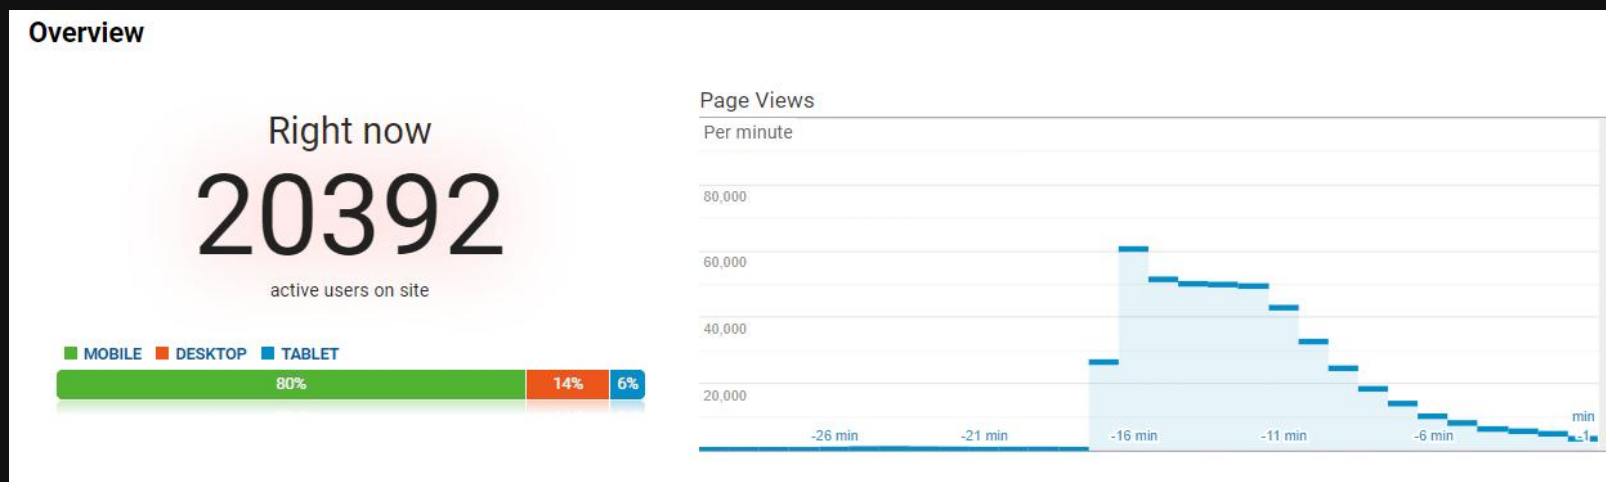
- Increase traffic above peak || decrease available resources
- Try to break your system
- Surface bottlenecks

SOAK TEST

- Extended test duration
- Watch behavior on ramp down as well as ramp up
 - Memory leaks
 - Disk space exhaustion (logs!)
 - Filled caches

SPIKE TEST: A STRESS TEST WITH QUICK RAMP-UP

- [Woot.com](#) at midnight
- TV ad "go online"
- System comes back online after downtime
- Everyone hits your API via on-the-hour cron jobs



Source: <https://twitter.com/troyhunt/status/1102312963401109504>

WHEN SHOULD YOU RUN A LOAD TEST?

- When your application performance may change
 - Adding or removing features
 - Refactoring
 - Infrastructure changes
- When your load profile may change
 - Initial app launch
 - Feature launch
 - Marketing pushes and promotions

HOW SHOULD I TEST?

**HOW SHOULD I TEST?
ACCURATELY.**

WHAT SHOULD I TEST?

- Flows (not just single endpoints)
- Frequently used
- Performance intensive
- Business critical

CONCURRENT REQUESTS != CONCURRENT USERS

- Think Time
- API client concurrency
- Caching (client-side or otherwise)

HOW **NOT** TO MODEL THINK TIME

- Ignore it
- Use a static amount
- Use a uniform distribution
(use a normal distribution instead)
- Assume you have one type of user

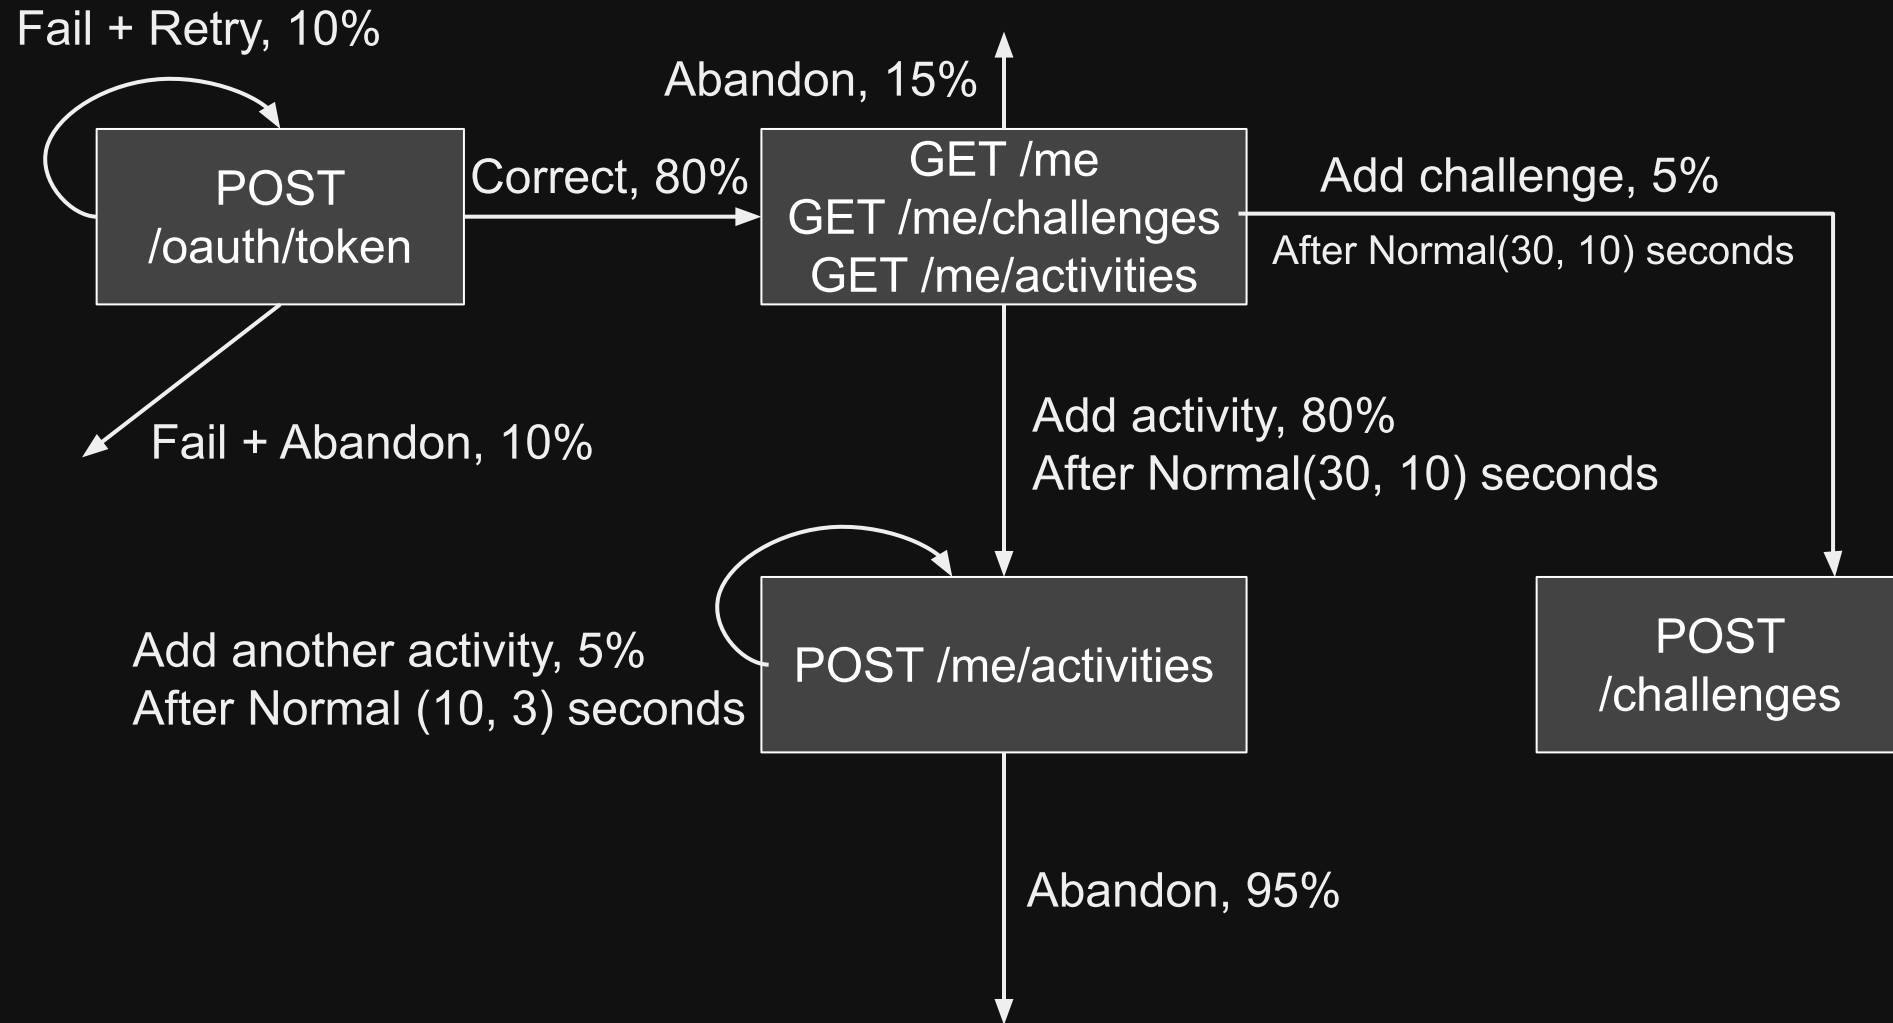
OVERSIMPLIFICATIONS TO AVOID

- No starting data in database
- No parameterization
- No abandonment at each step in the process
- No input errors

VARY YOUR TESTING

- High-load Case: more expensive endpoints get called more often
- Anticipated Case
- Low-load Case: validation failures + think time

SYSTEM UNDER TEST: CHALLENGR



LET'S SEE WHAT THAT LOOKS LIKE WITH K6

Yes, I should've used `const` instead of `let` everywhere.

IMPORTS

```
import http from "k6/http";  
import {check, fail, sleep} from "k6";  
import {Trend} from "k6/metrics";  
import {Normal} from [some gist URL];  
// Browserified AndreasMadsen/distributions
```

AUTH + FIXTURE DATA

```
// baseURL e.g. http://my-load-test-system.local/  
const [baseURL, clientId, clientSecret] =  
    open('./config.txt').split("\n"),  
    // start on the second line of the document, one per line  
    emails = open("./emails.csv").split("\n").slice(1),
```


PROBABILITIES AND INPUT SPECS

```
pCorrectCredentials = 0.8,  
pRetryAfterFailedCreds = 0.5,  
pAbandonAfterHomeLoad = 0.15,  
pAddChallenge = 0.05,  
pAddAnotherActivity = 0.05,  
pIncludeChallengeDuration = 0.5,  
pIncludeChallengeMileage = 0.5,  
// start with larger units for more accurate approximation  
// of what challenges look like  
challengeMinHalfHours = 1,  
challengeMaxHalfHours = 80,  
challengeMinTenMiles = 1,  
challengeMaxTenMiles = 20,  
activitySpeed = new Normal(15, 3),  
activityMinSeconds = 180,  
activityMaxSeconds = 10800,
```

THINK TIME DISTRIBUTIONS

```
challengeThinkTime = new Normal(30, 10),  
activityThinkTime = new Normal(30, 10),  
secondActivityThinkTime = new Normal(10, 3),
```

TRENDS: HOW LONG DID IT TAKE?

```
challengeListResponseTime  
    = new Trend("challenge_list_response_time"),  
activityListResponseTime  
    = new Trend("activity_list_response_time"),  
userProfileResponseTime  
    = new Trend("user_profile_response_time");
```

NOW THAT OUR SETUP IS DONE...

```
export default function () {  
    let isIncorrectLogin = Math.random() > pCorrectCredentials,  
        email = emails[getRandomInt(0, emails.length)];
```

LET'S MAKE AN HTTP CALL

```
let resLogin = http.post(baseUrl + "oauth/token", {  
  "client_id": clientId,  
  "client_secret": clientSecret,  
  "grant_type": "password",  
  "username": email,  
  "password": isIncorrectLogin ? "seekrit" : "secret",  
}, {  
  headers: {  
    "Content-Type": "application/x-www-form-urlencoded"  
  }  
})
```

MAKE SURE WE FAIL SUCCESSFULLY

```
if (isIncorrectLogin) {  
    check(resLogin, {  
        "invalid login caught": (res) => res.status === 401  
    }) || fail("no 401 on invalid login");  
  
    if (Math.random() > pRetryAfterFailedCreds) {  
        return; // abandon on incorrect login  
    }  
  
    // log in the correct way this time  
    resLogin = http.post(baseUrl + "oauth/token", {  
        "client_id": clientId,  
        // ...snip...  
    })  
}
```

MAKE SURE WE SUCCEED SUCCESSFULLY

```
check(resLogin, {  
  "login succeeded": (res) => res.status === 200  
    && typeof res.json().access_token !== "undefined",  
}) || fail("failed to log in");
```

MAKING SIMULTANEOUS REQUESTS

```
let params = {
  headers: {
    "Content-Type": "application/json",
    "Accept": "application/json",
    "Authorization": "Bearer " + resLogin.json().access_token
  }
}, makeGet = function (path) {
  return {method: "GET", url: baseUrl + path, params: params};
};

let homeScreenResponses = http.batch({
  "me": makeGet("api/me"),
  "challenges": makeGet("api/me/challenges"),
  "activities": makeGet("api/me/activities")
});
```


CHECKING SIMULTANEOUS RESPONSES

```
check(homeScreenResponses["me"],
      {"User profile loaded": (res) => res.json().email === email})
  || fail("user profile email did not match");
check(homeScreenResponses["challenges"],
      {"Challenges list loaded": (res) => res.status === 200})
  || fail("challenges list GET failed");
check(homeScreenResponses["activities"],
      {"Activities list loaded": (res) => res.status === 200})
  || fail("activities list GET failed");
```

TIMING SIMULTANEOUS RESPONSES

```
activityListResponseTime
    .add(homeScreenResponses["activities"].timings.duration);
challengeListResponseTime
    .add(homeScreenResponses["challenges"].timings.duration);
userProfileResponseTime
    .add(homeScreenResponses["me"].timings.duration);
```

DECIDE WHAT TO DO NEXT

```
let pNextAction = Math.random();  
if (pNextAction > (1 - pAbandonAfterHomeLoad)) {  
    return; // abandon here  
} else if (pNextAction >  
    (1 - pAbandonAfterHomeLoad - pAddChallenge)) {  
    // think time before creating challenge  
    sleep(fromDist(challengeThinkTime));
```

LET'S CREATE A CHALLENGE

```
let startMonth = getRandomInt(1, 3), endMonth = startMonth + getRandomInt(1, 2),
    challengeRes = http.post(baseUrl + "api/challenges", JSON.stringify({
      "name": "Test Challenge",
      "starts_at": "2020-0" + startMonth + "-01 00:00:00",
      "ends_at": "2020-" + (endMonth >= 10
        ? endMonth : ("0" + endMonth)) + "-01 00:00:00",
      "duration": Math.random() > pIncludeChallengeDuration ? null
        : secondsToTime(
            getRandomInt(challengeMinHalfHours, challengeMaxHalfHours) * 1800),
      "distance_miles": Math.random() > pIncludeChallengeMileage ? null
        : getRandomInt(challengeMinTenMiles, challengeMaxTenMiles) * 10
    }), params);
```

CHALLENGE ACCEPTED?

```
check(challengeRes, {"challenge was created":  
    (res) => res.status === 201 && res.json().id  
}) || fail("challenge create failed");  
  
let challengeListRes = http.get(baseUrl + "api/me/challenges", params);  
check(challengeListRes, {  
    "challenge is in user challenge list": (res) => {  
        let json = res.json();  
        for (let i = 0; i < json.created.length; i++)  
            if (json.created[i].id === challengeRes.json().id)  
                return true;  
        return false;  
    }  
}) || fail("challenge was not in user challenge list");
```

...OR WE CREATE ACTIVITIES THE SAME WAY

UNDERSTAND YOUR LOAD TEST TOOL

FOR EXAMPLE, **ARRIVAL RATE VS. LOOPING**

k6 is **working on it**...slowly...

AGGREGATE YOUR METRICS RESPONSIBLY

- ~~Average~~
- Median (~50th percentile)
- 90th, 95th, 99th percentile
- Standard Deviation
- Distribution of results
- Explain (don't discard) your outliers

KEEP IT REAL

- Use logs and analytics to determine your usage patterns
- Run your APM (e.g. New Relic) on your system under test
 - Better profiling info
 - Same performance drop as instrumenting production
- Is your infrastructure code? (e.g. Terraform, CloudFormation)
 - Easier to copy environments
 - Cheaper to set up an environment for an hour to run a load test
- Decide whether testing from near your env is accurate enough
- Test autoscaling and/or load-shedding facilities

LET'S RUN ANOTHER TEST!

WARNING: TRICKY BOTTLENECKS AHEAD

- Just because a request is expensive
doesn't mean it's the biggest source of load
- As a system reaches capacity
you'll see nonlinear performance degradation

BOTTLENECKS: WEB SERVER + DATABASE

- Web workers (e.g. FPM)/Apache processes
- DB Connections
- CPU + RAM utilization
- Network utilization
- Disk utilization (I/O or space)

BOTTLENECKS: LOAD BALANCER

- Network utilization/warmup
- Connection count

BOTTLENECKS: EXTERNAL SERVICES

- Rate limits (natural or artificial)
- Latency
- Network egress

BOTTLENECKS: QUEUES

- Per-job spin-up latency
- Worker count
- CPU + RAM utilization
 - Workers
 - Broker
- Queue depth

BOTTLENECKS: CACHES

- Thundering herd
- Churning due to cache evictions

LET'S FIX SOME BOTTLENECKS...

BONUS: MORE TOOLS

- [Artillery.io](#) (Node)
 - Simple stuff in Yaml
 - Can switch to JS (with npm)
- [Molotov](#) (by Mozilla, in Python)
 - Uses async IO via coroutines
- [Locust](#) (Python)
 - Can be run clustered
- [Siege](#)
 - [GitLab Large Staging Collider](#)
- [Apache JMeter](#) (Java)
- [Gatling](#) (Java)
 - Tests in Scala...
 - ...or use the recorder
- [ab](#)
- [httperf](#)
- [Wrk2](#) (C)
 - Scriptable via Lua

WHAT WE LEARNED

- What types of tests exist, and when you should use them
- **How to match load tests with (anticipated) reality**
- What a real performance test script looks like in K6
- How to analyze results during and after your test

FURTHER READING

- [Performance Testing Guidance for Web Applications](#) (from Microsoft)
- [Blazemeter Blog](#) - solid info on load testing topics
- [ian.im/loadarch](#) - an article version of this talk (php[architect] sub req'd)
- [test-api.loadimpact.com](#) - an API to load test against from the k6 folks

THANKS! QUESTIONS?

- ian.im/loadfoo20 - these slides
- github.com/iansltx/challenger - this code
- twitter.com/iansltx - me
- github.com/iansltx - my code
- Please leave feedback; thanks :)