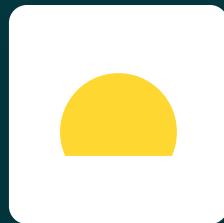




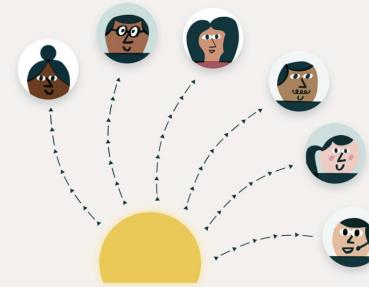
Redis is not just a cache

CONFOO MONTREAL 2020



Reach users wherever they are with a single API

Sunshine Conversations' universal messaging API lets you quickly and easily communicate with users on all the major chat platforms.



WHAT IS REDIS?

what is redis

“Redis is an open source (BSD licensed),
in-memory data structure store, used as a
database, cache and message broker.”

(In the creator's own words)

“Redis is an open source (BSD licensed),
in-memory data structure store, used as a
database, **cache** and message broker.”

(In the creator's own words)



“Redis is an open source (BSD licensed),
in-memory **data structure store**, used as a
database, cache and message broker.”

(In the creator's own words)

“Redis is an open source (BSD licensed),
in-memory data structure store, used as a
database, cache and message broker.”

(In the creator's own words)

in-memory does not mean ephemeral

redis also persists to disk

- full snapshot format (RDB)
- append only file format (AOF) “binlog”

in-memory does not mean ephemeral

configure your durability needs

it's a tradeoff with throughput

fsync every second

no fsync

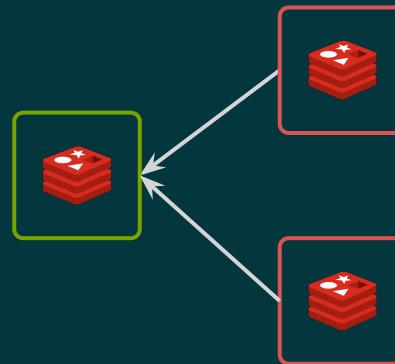
fsync every command

in-memory does not mean ephemeral

redis is also highly available (HA)

supports replication

supports sharding (aka clustering)



the commands

alphabet soup

SET key value [expiration EX seconds|PX milliseconds] [NX|XX]

Set the string value of a key

BLPOP key [key ...] timeout

Remove and get the first element in a list, or block until one is available

BRPOP key [key ...] timeout

Remove and get the last element in a list, or block until one is available

BRPOPLPUSH source destination timeout

Pop a value from a list, push it to another list and return it; or block until one is available

BZPOPMIN key [key ...] timeout

Remove and return the member with the lowest score from one or more sorted sets, or block until one is available

BZPOPMAX key [key ...] timeout

Remove and return the member with the highest score from one or more sorted sets, or block until one is available

HMSET

HMSET key field value [field value ...]

Available since 2.0.0.

Time complexity: O(N) where N is the number of fields being set.

Sets the specified fields to their respective values in the hash stored at key. This command overwrites any specified fields already existing in the hash. If key does not exist, a new key holding a hash is created.

H M SET

H is for Hash, the type of data structure we're setting

HMSET key field value [field value ...]

Available since 2.0.0.

Time complexity: O(N) where N is the number of fields being set.

Sets the specified fields to their respective values in the hash stored at key. This command overwrites any specified fields already existing in the hash. If key does not exist, a new key holding a hash is created.

H M SET

H is for Hash, the type of data structure we're setting

M is for multi, because we're setting multiple keys at once

HMSET key field value [field value ...]

Available since 2.0.0.

Time complexity: O(N) where N is the number of fields being set.

Sets the specified fields to their respective values in the hash stored at key. This command overwrites any specified fields already existing in the hash. If key does not exist, a new key holding a hash is created.

H M SET

H is for Hash, the type of data structure we're setting

M is for multi, because we're setting multiple keys at once

SET is for set... that's what we're doing

HMSET key field value [field value ...]

Available since 2.0.0.

Time complexity: O(N) where N is the number of fields being set.

Sets the specified fields to their respective values in the hash stored at key. This command overwrites any specified fields already existing in the hash. If key does not exist, a new key holding a hash is created.

H M SET

H is for Hash, the type of data structure we're setting

M is for multi, because we're setting multiple keys at once

SET is for set... that's what we're doing

HMSET key field value [field value ...]

Available since 2.0.0.

Time complexity: O(N) where N is the number of fields being set.

Sets the specified fields to their respective values in the hash stored at key. This command overwrites any specified fields already existing in the hash. If key does not exist, a new key holding a hash is created.

```
> hmset myhash name andrew job developer  
OK
```

H M SET

H is for Hash, the type of data structure we're setting

M is for multi, because we're setting multiple keys at once

SET is for set... that's what we're doing

HMSET key field value [field value ...]

Available since 2.0.0.

Time complexity: O(N) where N is the number of fields being set.

Sets the specified fields to their respective values in the hash stored at key. This command overwrites any specified fields already existing in the hash. If key does not exist, a new key holding a hash is created.

```
> hmset myhash name andrew job developer
OK
> hgetall myhash
1) "name"
2) "andrew"
3) "job"
4) "developer"
```

REDIS COMMANDS

There are command prefixes for each data structure

key	= no prefix
hash	= H
list	= L or R (left or right)
sets	= S
sorted sets	= Z
streams	= X

And a few other prefixes, postfixes and options

M	= get or set multiple
X	= postfix, only proceed if key exists
B	= blocking, if there's nothing to get wait indefinitely until there is
XX	= only proceed if key exists
NX	= only proceed if key <i>does not</i> exist
PX	= set an expiration on this key

alphabet soup

BRPOPLPUSH ??👉

SET key value [expiration EX seconds|PX milliseconds] [NX|XX]

Set the string value of a key

BLPOP key [key ...] timeout

Remove and get the first element in a list, or block until one is available

BRPOP key [key ...] timeout

Remove and get the last element in a list, or block until one is available

BRPOPLPUSH source destination timeout

Pop a value from a list, push it to another list and return it; or block until one is available

BZPOPMIN key [key ...] timeout

Remove and return the member with the lowest score from one or more sorted sets, or block until one is available

BZPOPMAX key [key ...] timeout

Remove and return the member with the highest score from one or more sorted sets, or block until one is available

BRPOPLPUSH

B RPOP LPUSH

B is for blocking

B RPOP LPUSH

B is for blocking

RPOP pop from the right side of one list

B RPOP LPUSH

B is for blocking

RPOP pop from the right side of one list

LPUSH push the item to left side of another list

B RPOP LPUSH

B is for blocking

RPOP pop from the right side of one list

LPUSH push the item to left side of another list

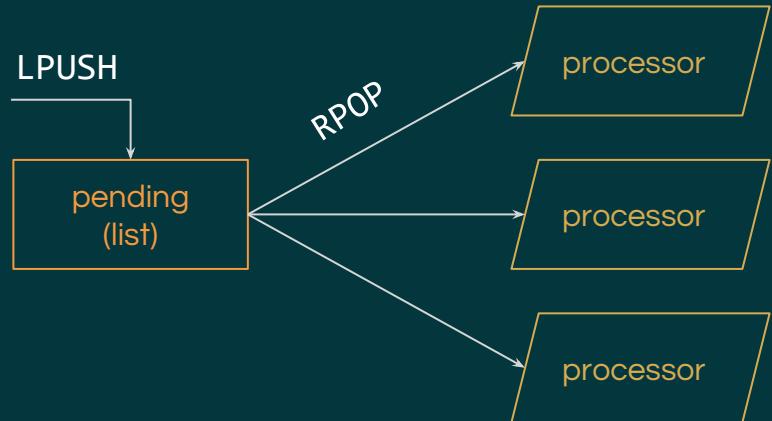
(do it atomically)

B RPOP LPUSH

B is for blocking

RPOP pop from the right side of one list

LPUSH push the item to left side of another list

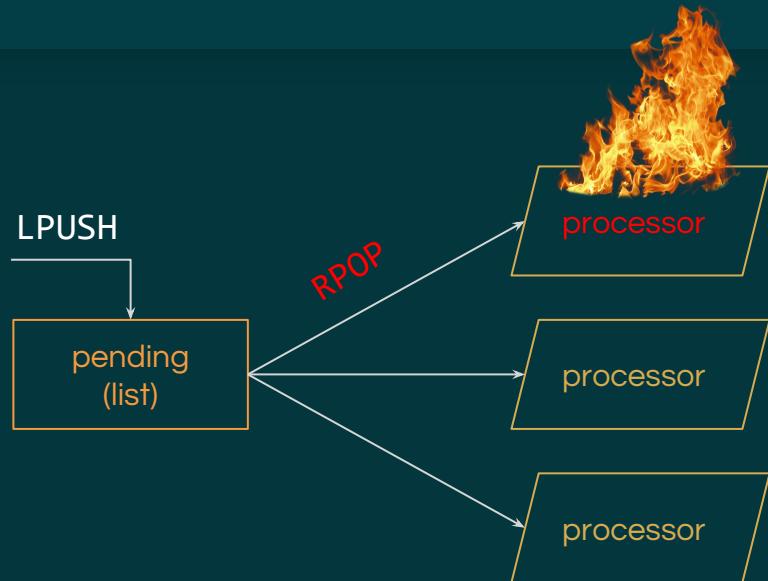


B RPOP LPUSH

B is for blocking

RPOP pop from the right side of one list

LPUSH push the item to left side of another list

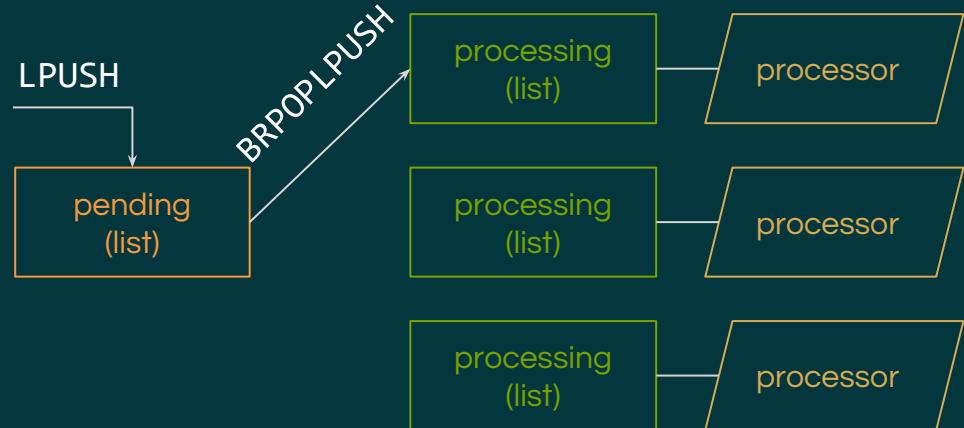


B RPOP LPUSH

B is for blocking

RPOP pop from the right side of one list

LPUSH push the item to left side of another list

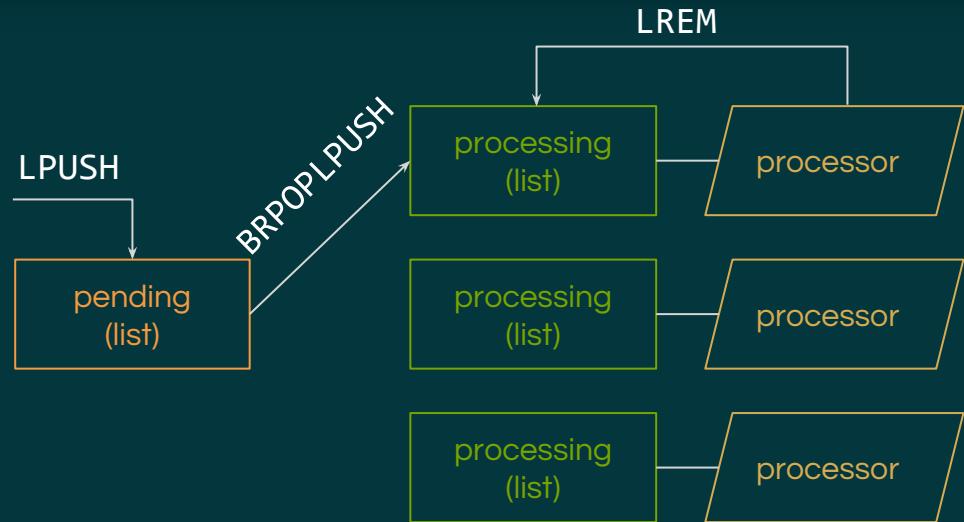


B RPOP LPUSH

B is for blocking

RPOP pop from the right side of one list

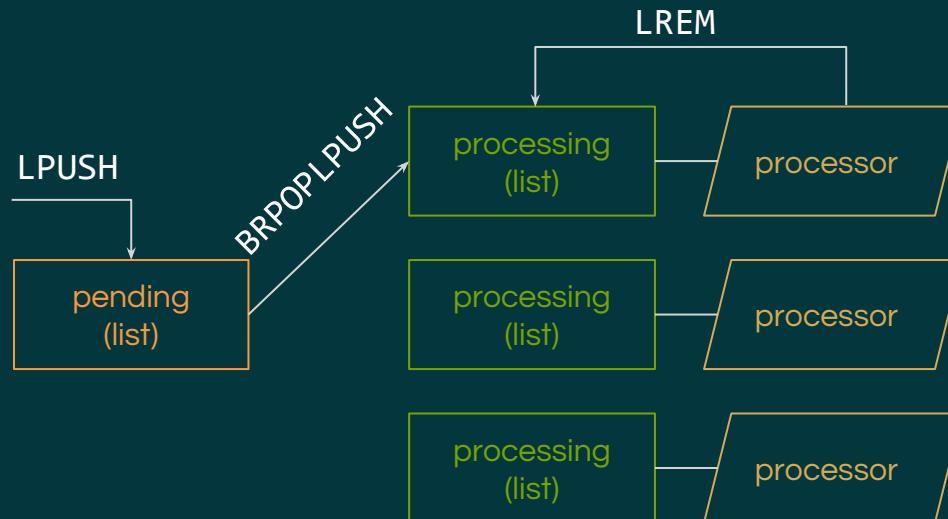
LPUSH push the item to left side of another list



REDIS COMMANDS

```
127.0.0.1:6379> brpoplpush src dest 0
```

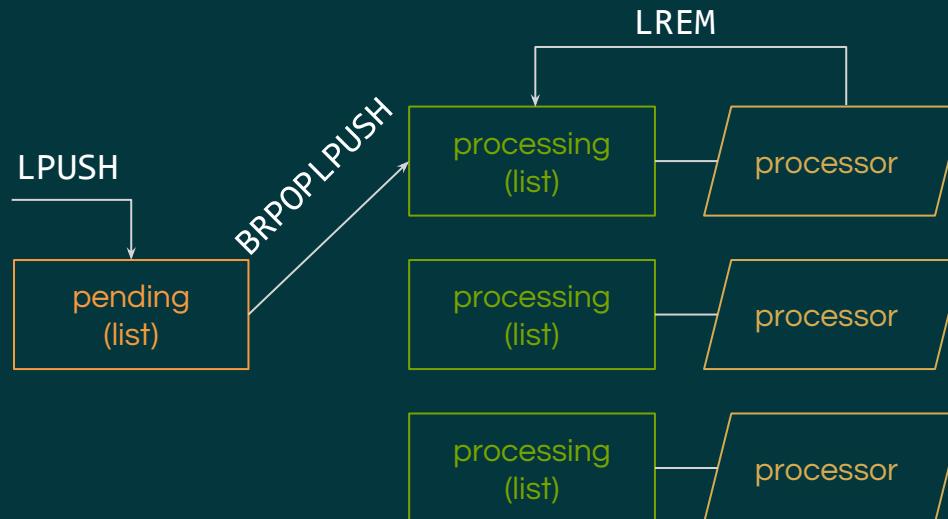
```
127.0.0.1:6379> rpush src he█[value ...]
```



REDIS COMMANDS

```
127.0.0.1:6379> brpoplpush src dest 0  
"hello"  
(4.90s)  
127.0.0.1:6379>
```

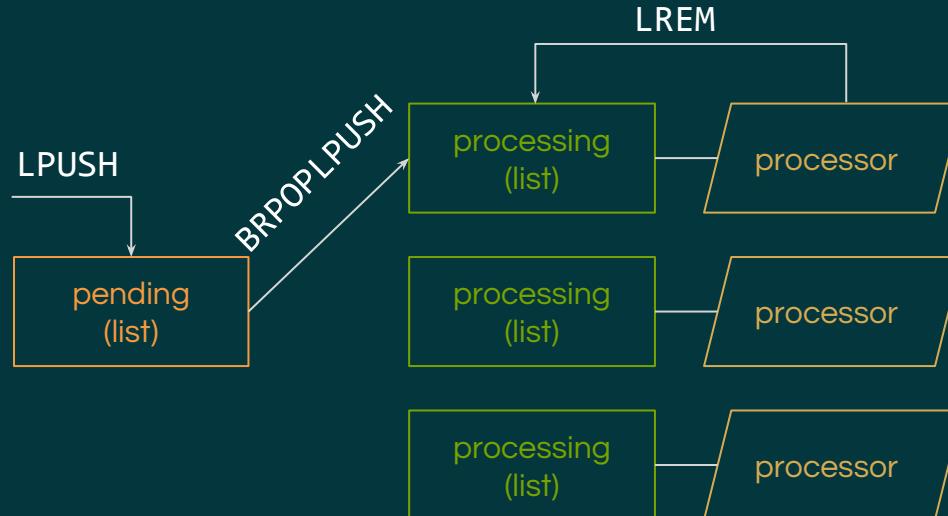
```
127.0.0.1:6379> rpush src hello  
(integer) 1  
127.0.0.1:6379> ■
```



REDIS COMMANDS

```
127.0.0.1:6379> brpoplpush src dest 0  
"hello"  
(4.90s)  
127.0.0.1:6379>
```

```
127.0.0.1:6379> rpush src hello  
(integer) 1  
127.0.0.1:6379> lrange dest 0 -1  
1) "hello"  
127.0.0.1:6379> █
```



alphabet soup

Related commands

- [BLPOP](#)
- [BRPOP](#)
- [**BRPOPLPUSH**](#)
- [LINDEX](#)
- [LINSERT](#)
- [LLEN](#)
- [LPOP](#)
- [LPUSH](#)
- [LPUSHX](#)
- [LRANGE](#)
- [LREM](#)
- [LSET](#)
- [LTRIM](#)
- [RPOP](#)
- [RPOPLPUSH](#)
- [RPUSH](#)
- [RPUSHX](#)

alphabet
soup is
starting to
make sense

Related commands

- [BLPOP](#)
- [BRPOP](#)
- [**BRPOPLPUSH**](#)
- [LINDEX](#)
- [LINSERT](#)
- [LLEN](#)
- [LPOP](#)
- [LPUSH](#)
- [LPUSHX](#)
- [LRANGE](#)
- [LREM](#)
- [LSET](#)
- [LTRIM](#)
- [RPOP](#)
- [RPOPLPUSH](#)
- [RPUSH](#)
- [RPUSHX](#)

THE MEAT

the meat

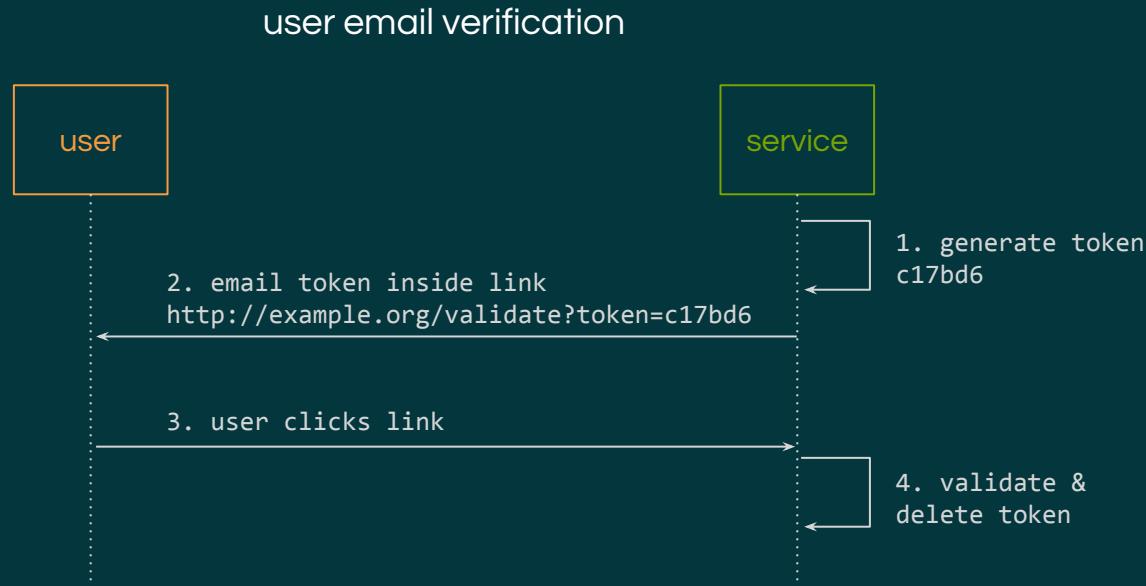


THE MEAT

- 1 ONE TIME TOKENS
- 2 LOCKS
- 3 RATE LIMITS

1. One time tokens

ONE TIME TOKENS



ONE TIME TOKENS

issue a new token:

```
> setex code:c17bd6 3600000 '{"act_id": "1234", "email": "user@example.org"}'  
OK
```

ONE TIME TOKENS

issue a new token:

```
> setex code:c17bd6 3600000 '{"act_id": "1234", "email": "user@example.org"}'  
OK
```

this is the key

the colon is just a delimiter to keep things
organized

ONE TIME TOKENS

issue a new token:

```
> setex code:c17bd6 3600000 '{"act_id": "1234", "email": "user@example.org"}'  
OK
```

the token itself



ONE TIME TOKENS

issue a new token:

```
> setex code:c17bd6 3600000 '{"act_id": "1234", "email": "user@example.org"}'  
OK
```

the token itself
ms before it expires (1 hour)

ONE TIME TOKENS

issue a new token:

```
> setex code:c17bd6 3600000 '{"act_id": "1234", "email": "user@example.org"}'  
OK
```

the token itself

ms before it expires (1 hour)

the value of the key

ONE TIME TOKENS

verify email is sent to the user
and the user clicks the link...



https://yoursite.com/verify_email?token=c17bd6

ONE TIME TOKENS

now we validate the token

```
> get code:c17bd6
"{"act_id": "1234", "email": "user@example.org"}"
```

and consume it

```
> del code:c17bd6
(integer) 1
```

we can do both **get** and **del** in one shot

you begin a redis transaction with **multi**
and you end it with **exec**

commands in between are
executed as an atomic transaction

```
> multi  
OK  
  
> any number of commands  
  
> exec  
array of command results
```



ONE TIME TOKENS

validate & consume the token

```
> multi
OK
> get code:c17bd6
QUEUED
> del code:c17bd6
QUEUED
> exec
1) "{\"act_id\": \"1234\", \"email\": \"user@example.org\"}"
2) (integer) 1
```

ONE TIME TOKENS

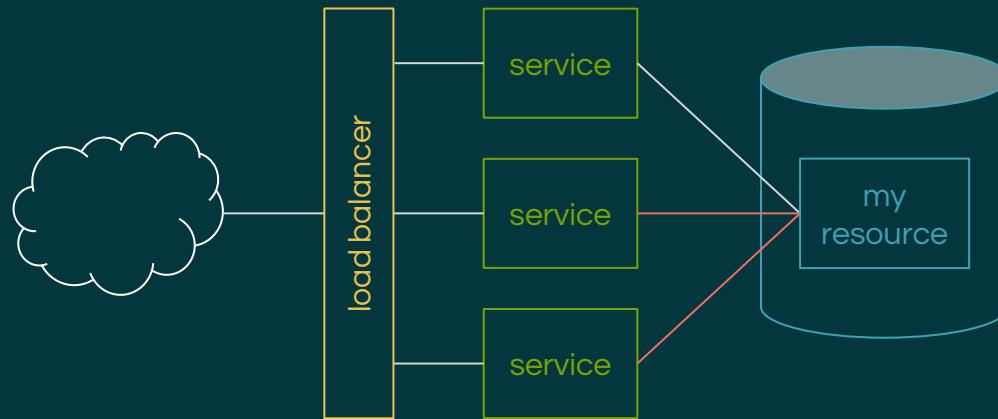
reject an invalid or consumed code

```
> multi
OK
> get code:invalid
QUEUED
> del code:invalid
QUEUED
> exec
1) (nil)
2) (integer) 0
```

2. Locks

LOCKS

why locks?



LOCKS

a redis key can be a lock

```
> set myresource ef1f63c NX PX 10000
OK
```

LOCKS

a redis key can be a lock

```
> set myresource ef1f63c NX PX 10000
OK
```

the resource id

LOCKS

a redis key can be a lock

```
> set myresource ef1f63c NX PX 10000
OK
```

the resource id
the key value

a redis key can be a lock

```
> set myresource ef1f63c NX PX 10000  
OK
```

the resource id

the key value

NX = don't overwrite any existing key

LOCKS

a redis key can be a lock

```
> set myresource ef1f63c NX PX 10000
OK
```

the resource id

the key value

NX = don't overwrite any existing key

PX = expire this key in 10 seconds

LOCKS

OK means the lock was acquired

```
> set myresource ef1f63c NX PX 10000
OK
```

(nil) means the key already exists; try again later

```
> set myresource ef1f63c NX PX 10000
(nil)
```

LOCKS

OK means the lock was acquired

```
> set myresource ef1f63c NX PX 10000
OK
```

(nil) means the key already exists; try again later

```
> set myresource ef1f63c NX PX 10000
(nil)
```

LOCKS

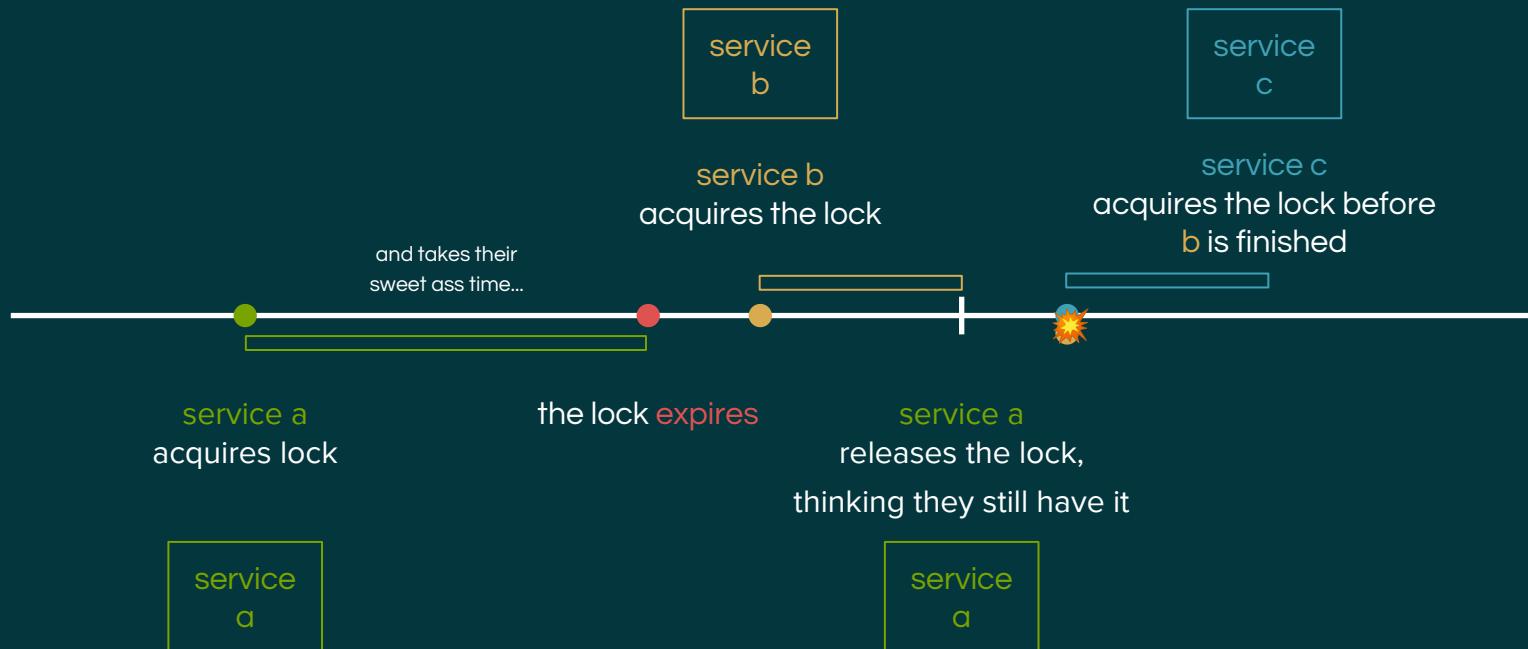
release the lock by deleting the key

```
> del myresource  
(integer) 1
```

Is it really that
simple?

Unfortunately
not

LOCKS



release the lock by deleting the key

```
> del myresource  
(integer) 1
```

we need some way to check that we have the lock before releasing it

and we somehow need to do this part transactionally...

```
lock_value = crypto.random()  
  
> set myresource lock_value NX PX 10000  
OK  
  
do_my_stuff()  
  
> get myresource  
"ef1f63c"  
  
if (lock_value === "ef1f63c") {  
  > del myresource  
}
```

Lua scripting!

EVAL and **EVALSHA** commands allow you to run some logic
remotely in redis

Redis will execute your lua script atomically

our safe release lua script looks like this:

```
if redis.call('get', KEYS[1]) == ARGV[1]
then return redis.call('del', KEYS[1])
else return 0 end
```

our safe release lua script looks like this:

```
if redis.call('get', KEYS[1]) == ARGV[1]
then return redis.call('del', KEYS[1])
else return 0 end
```

lua scripts use `redis.call()` to invoke redis commands

our safe release lua script looks like this:

```
if redis.call('get', KEYS[1]) == ARGV[1]
then return redis.call('del', KEYS[1])
else return 0 end
```

arguments are passed to the script as key value pairs

we pass the resource key as KEYS[1]
and the lock value ARGV[1]

our safe release lua script looks like this:

```
if redis.call('get', KEYS[1]) == ARGV[1]
then return redis.call('del', KEYS[1])
else return 0 end
```

arguments are passed to the script as key value pairs

we pass the resource key as KEYS[1]
and the lock value ARGV[1]

(and yes... Lua has 1-based indexes 😊)

our safe release lua script looks like this:

```
if redis.call('get', KEYS[1]) == ARGV[1]
then return redis.call('del', KEYS[1])
else return 0 end
```

LOCKS

acquire the lock (after generating a lock value)

```
> set myresource ef1f63c NX PX 10000
OK
```

release the lock (safely)

```
> eval "if redis.call('get', KEYS[1]) == ARGV[1] then return redis.call('del',
KEYS[1]) else return 0 end" 1 myresource , ef1f63c
(integer) 1
```

LOCKS

acquire the lock (after generating a lock value)

```
> set myresource ef1f63c NX PX 10000  
OK
```

release the lock (safely) the number of keys I'm passing

```
> eval "if redis.call('get', KEYS[1]) == ARGV[1] then return redis.call('del',  
KEYS[1]) else return 0 end" 1 myresource , ef1f63c  
(integer) 1
```

with **SCRIPT LOAD** and **EVALSHA** you don't have send your entire script payload every time

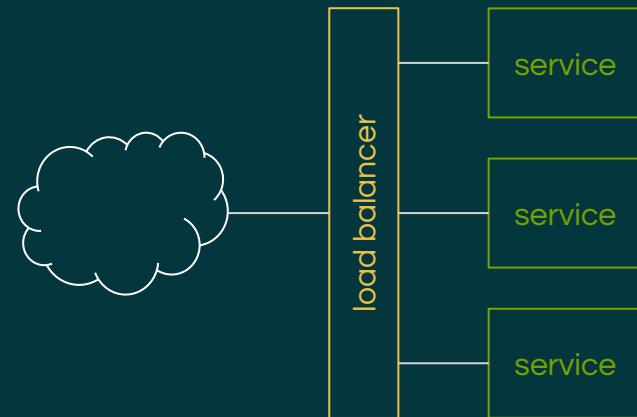
```
> script load "if redis.call('get', KEYS[1]) == ARGV[1] then return  
redis.call('del', KEYS[1]) else return 0 end"  
"e9f69f2beb755be68b5e456ee2ce9aadfbcc4ebf4"  
  
> evalsha e9f69f2beb755be68b5e456ee2ce9aadfbcc4ebf4 1 myresource , ef1f63c  
(integer) 1
```

3. Rate limits

RATE LIMITS

rate limits protect your service from traffic spikes

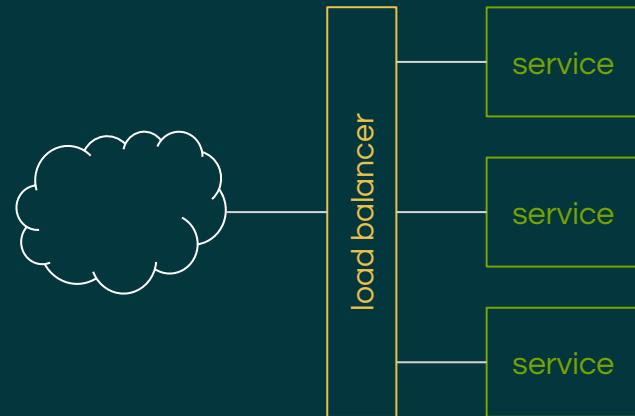
you'd typically them in your load balancer or web server,
e.g. ELB or NGINX



but perhaps your rate limit criteria is complex

e.g. what about a GraphQL **query complexity** rate limit

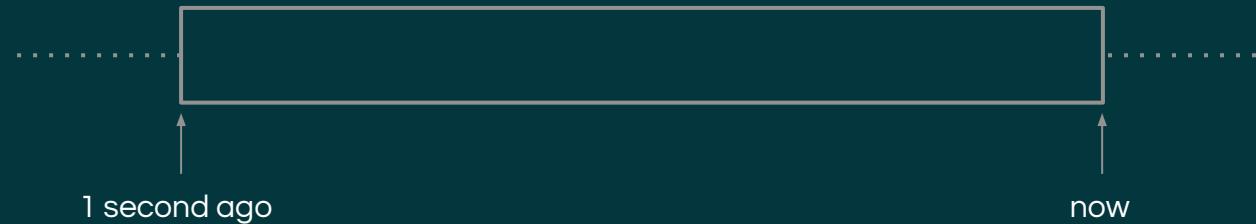
redis can help



RATE LIMITS

let's impose a limit of **5 requests per second**
starting simple, consider a **rolling window model**

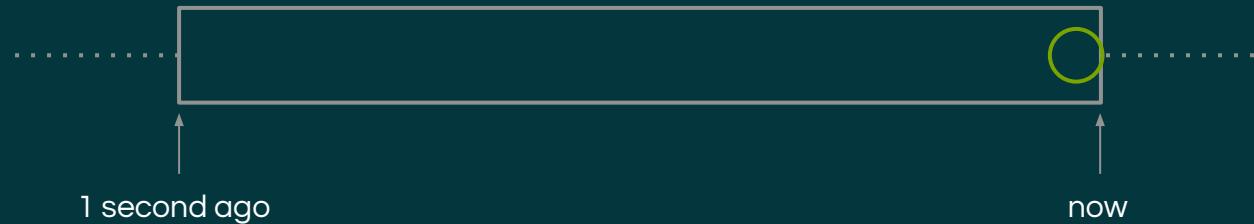
requests made during last second: **0** (limit 5)



RATE LIMITS

new requests come in at the front of the window

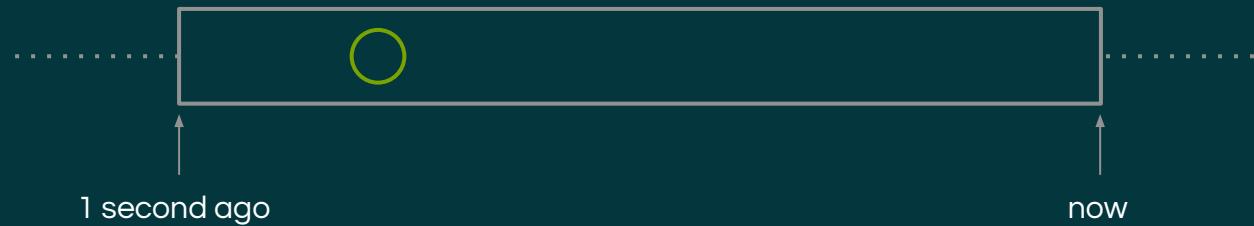
requests made during last second: **1** (limit 5)



RATE LIMITS

and as time goes on they drift to the back

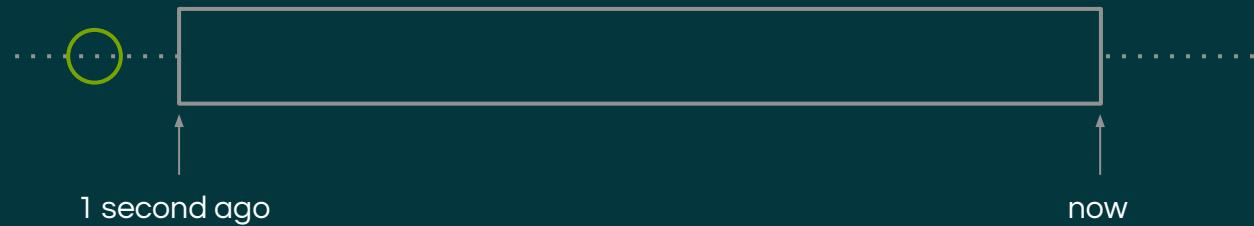
requests made during last second: **1** (limit 5)



RATE LIMITS

and out of the window

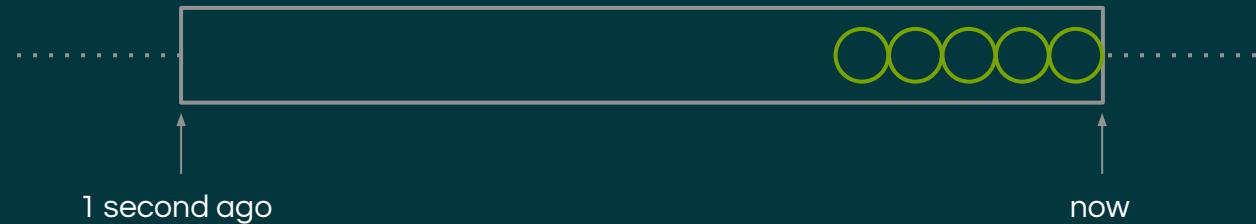
requests made during last second: **0** (limit 5)



RATE LIMITS

we can accept a burst of many requests at once

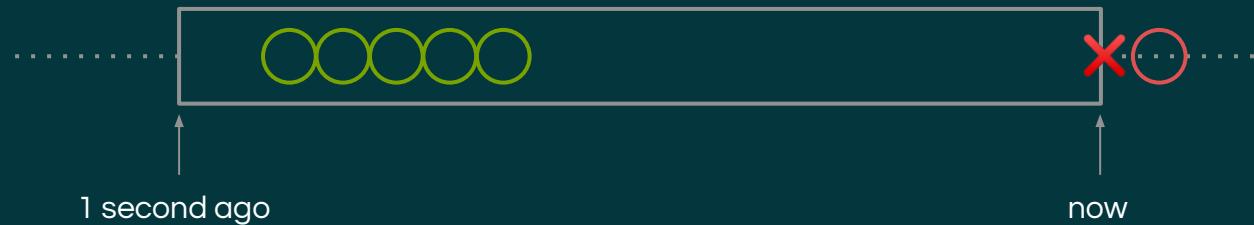
requests made during last second: **5** (limit 5)



RATE LIMITS

but some time needs to pass before we accept any more

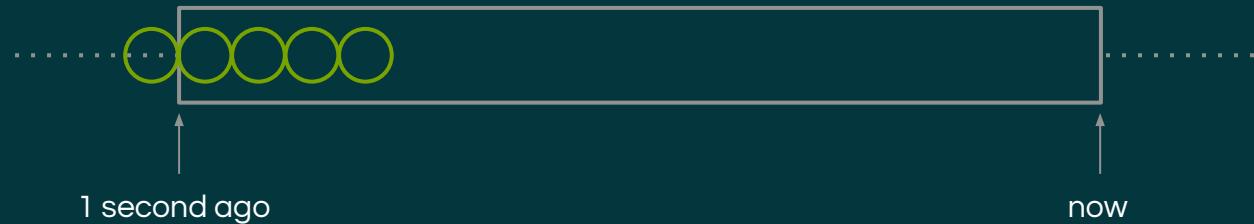
requests made during last second: **5** (limit 5)



RATE LIMITS

but some time needs to pass before we accept any more

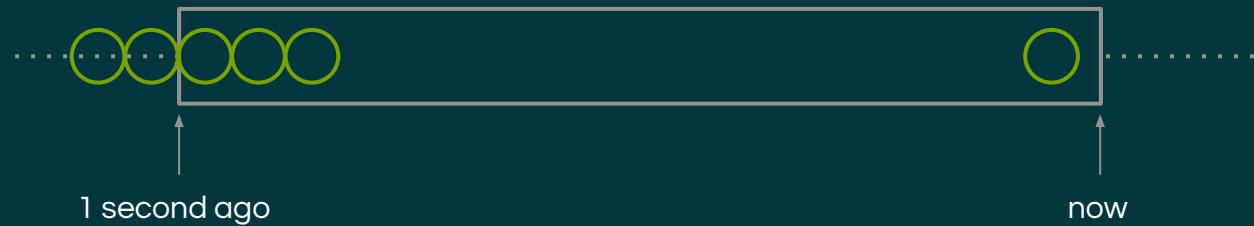
requests made during last second: **4** (limit 5)



RATE LIMITS

but some time needs to pass before we accept any more

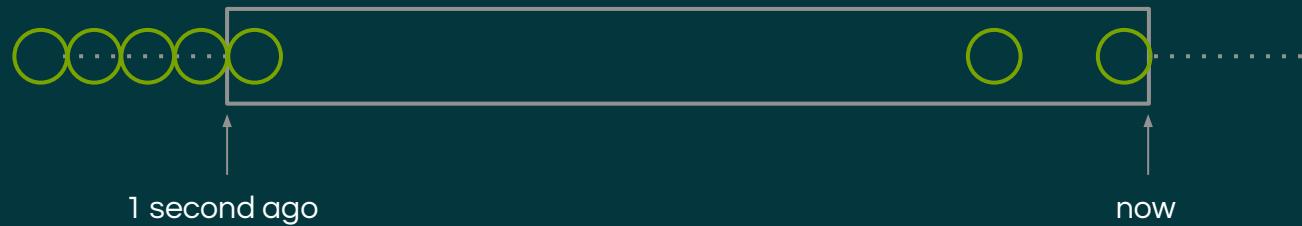
requests made during last second: **4** (limit 5)



RATE LIMITS

but some time needs to pass before we accept any more

requests made during last second: **3** (limit 5)



Let's make a rolling window with redis

we'll use lua



Let's make a rolling window with redis

we'll use lua
and sorted sets



Sorted Sets 101

Adding members (with scores)

```
> zadd leaderboard 1000000 parzival  
(integer) 1  
> zadd leaderboard 900000 art3mis  
(integer) 1
```

zrevrange: list members from start (**0**) to finish (**-1**), in reverse order

```
> zrevrange leaderboard 0 -1 withscores
1) "parzival"
2) "1000000"
3) "art3mis"
4) "900000"
5) "ioi987427"
6) "0"
7) "ioi677427"
8) "0"
9) "ioi423417"
10) "0"
```

Removing individual members

```
> zrem leaderboard ioi987427
(integer) 1
> zrevrange leaderboard 0 -1 withscores
1) "parzival"
2) "1000000"
3) "art3mis"
4) "900000"
5) "ioi677427"
6) "0"
7) "ioi423417"
8) "0"
```

RATE LIMITS

Removing a range of members

remove all members whose score is between `-inf` and `0`, inclusive:

```
> zremrangebyscore leaderboard -inf 0  
(integer) 2
```

```
> zrevrange leaderboard 0 -1 withscores  
1) "parzival"  
2) "1000000"  
3) "art3mis"  
4) "900000"
```

Tell me how many members there are (cardinality)

```
> zcard leaderboard  
(integer) 2
```

Now that you know **ZADD**,
ZCARD and
ZREMRANGEBYSCORE,
let's look at a Lua script

RATE LIMITS

```
local rateLimit = 5

local rateKey = KEYS[1]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    return 1;
end
```

👉 Much inspiration was had from Sahil Jadon's article *Rate Limiting using Redis Lists and Sorted Sets*
<https://medium.com/@sahiljadon/rate-limiting-using-redis-lists-and-sorted-sets-9b42bc192222>

RATE LIMITS

```
{  
    local rateLimit = 5  
    local rateKey = KEYS[1]  
    local ts = tonumber(ARGV[1])  
}  
  
redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)  
  
if redis.call('ZCARD', rateKey) < rateLimit then  
    redis.call('ZADD', rateKey, ts, ts)  
    return 0;  
else  
    return 1;  
end
```

Set the rate limit, receive args
Timestamp (ts) is passed in as an arg

RATE LIMITS

```
local rateLimit = 5

local rateKey = KEYS[1]
local ts = tonumber(ARGV[1])

{ redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000) }

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    return 1;
end
```

remove any members that are older than our rolling window



RATE LIMITS

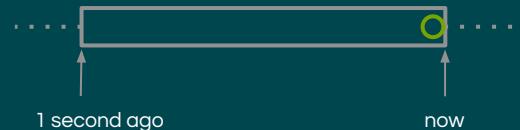
```
local rateLimit = 5

local rateKey = KEYS[1]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)

{
  if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
  else
    return 1;
end}
```

if there's room in the window,
accept the request



RATE LIMITS

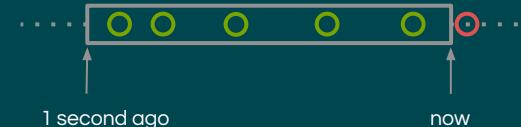
```
local rateLimit = 5

local rateKey = KEYS[1]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    return 1;
end
```

otherwise, reject the request and
return **1** to indicate



Let's test it

RATE LIMITS

eval the script 10 times per second
with a rate limit of 5, half get accepted

```
const script = fs.readFileSync('rolling_simple.lua', 'utf8');
const rc = redis.createClient();

setInterval(() => {
    rc.eval(script, 1, 'myresource', Date.now()).then(console.log);
}, 100);
```

```
$ node rolling_simple.js
```

```
0  
0  
0  
0  
0  
1  
1  
1  
1  
1  
1  
0  
0  
0  
0  
0
```

What about retry advice?

when returning a 429 HTTP rate limit error it's courteous to include a `Retry-After` header telling the client how long to wait

HTTP/1.1 429 Too Many Requests

`Content-Type: text/html`

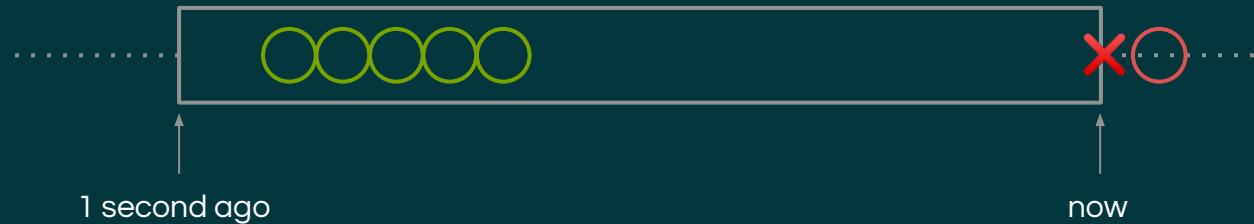
`Retry-After: 3600`

Let's compute a simple
Retry-After with redis

RATE LIMITS

remember our rolling window?

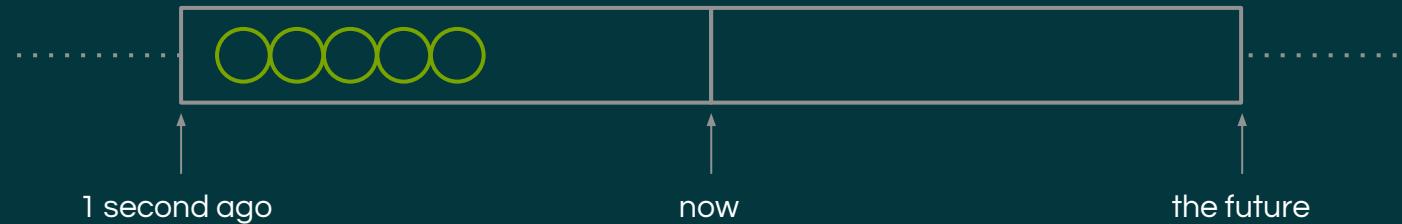
requests made during last second: **5** (limit 5)



RATE LIMITS

remember our rolling window?
let's add a *future* window

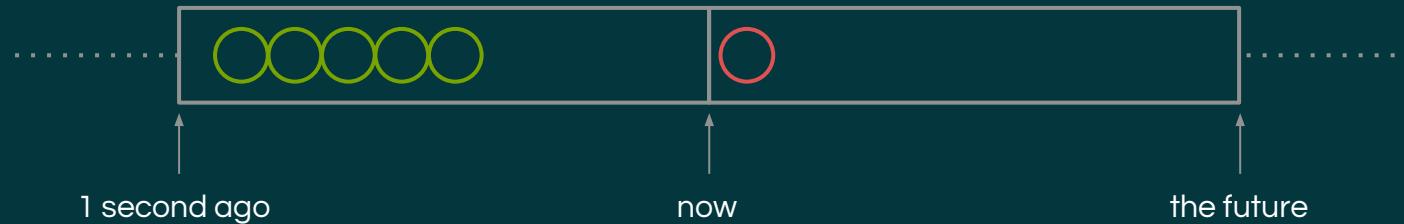
requests made during last second: **5** (limit 5)



RATE LIMITS

when request rates exceed our limit, add to our future window

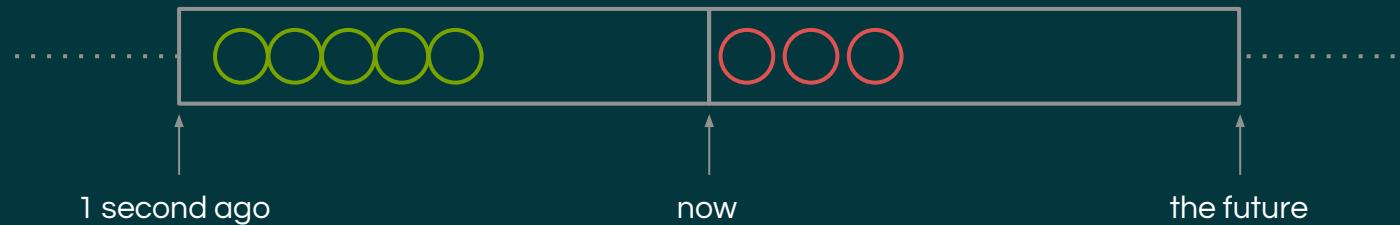
requests made during last second: **5** (limit 5)



RATE LIMITS

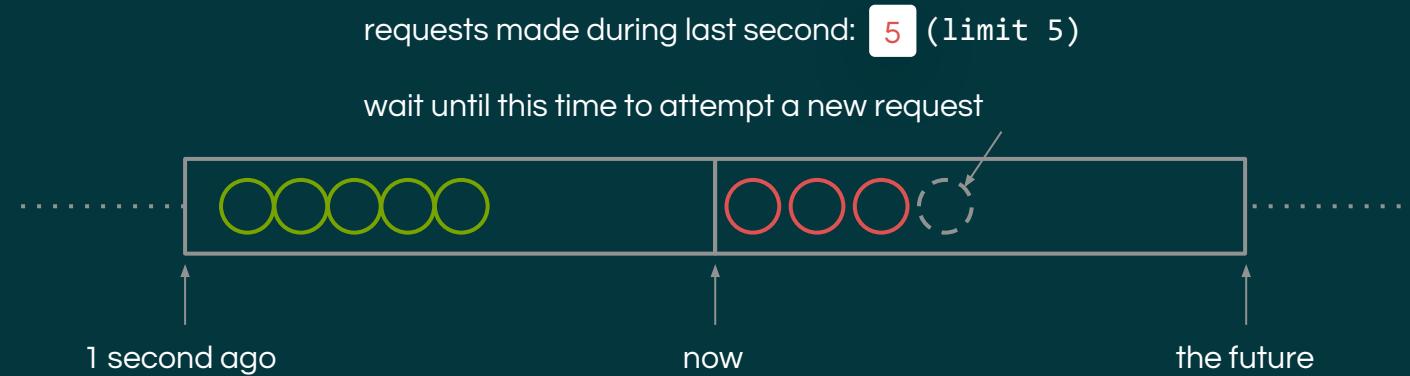
when request rates exceed our limit, add to our future window

requests made during last second: **5** (limit 5)



RATE LIMITS

the last value in our future window is the retry advice

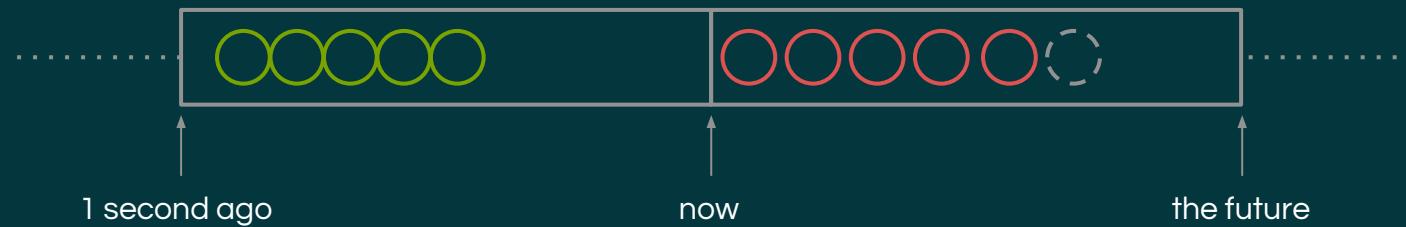


RATE LIMITS

as the rate limit continues to be exceeded the Retry-After will continue to grow

requests made during last second: **5** (limit 5)

wait until this time to attempt a new request

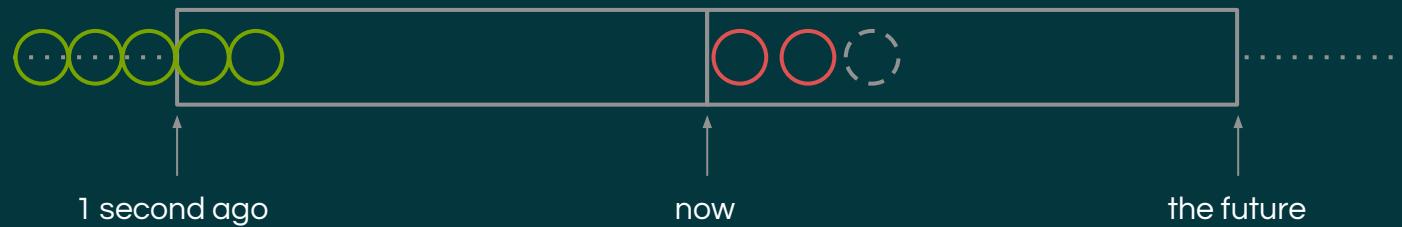


RATE LIMITS

and as time goes on our window moves leftward

requests made during last second: **2** (limit 5)

wait until this time to attempt a new request



RATE LIMITS

```
local rateLimit = 5

local rateKey = KEYS[1]
local retryKey = KEYS[2]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)
redis.call('ZREMRANGEBYSCORE', retryKey, '-inf', ts)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    local interval = 1000 / rateLimit

    local next = ts + interval;
    local last = redis.call('ZRANGE', retryKey, -1, -1, 'WITHSCORES')

    if type(last) == 'table' and #last > 0 then
        next = tonumber(last[#last]) + interval
    end

    redis.call('ZADD', retryKey, next, next);

    return next - ts;
end
```

RATE LIMITS

```
local rateLimit = 5
local rateKey = KEYS[1]
local retryKey = KEYS[2]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)
redis.call('ZREMRANGEBYSCORE', retryKey, '-inf', ts)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    local interval = 1000 / rateLimit

    local next = ts + interval;
    local last = redis.call('ZRANGE', retryKey, -1, -1, 'WITHSCORES')

    if type(last) == 'table' and #last > 0 then
        next = tonumber(last[#last]) + interval
    end

    redis.call('ZADD', retryKey, next, next);

    return next - ts;
end
```

this should all look familiar,
except we have two sorted sets
now, `rateKey` and `retryKey`

RATE LIMITS

```
local rateLimit = 5

local rateKey = KEYS[1]
local retryKey = KEYS[2]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)
redis.call('ZREMRANGEBYSCORE', retryKey, '-inf', ts)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    local interval = 1000 / rateLimit
    local next = ts + interval;
    local last = redis.call('ZRANGE', retryKey, -1, -1, 'WITHSCORES')
    if type(last) == 'table' and #last > 0 then
        next = tonumber(last[#last]) + interval
    end
    redis.call('ZADD', retryKey, next, next);
    return next - ts;
end
```

rate limit reached!

find the last member of the retry sorted set and add our interval to it



RATE LIMITS

```
local rateLimit = 5

local rateKey = KEYS[1]
local retryKey = KEYS[2]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)
redis.call('ZREMRANGEBYSCORE', retryKey, '-inf', ts)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    local interval = 1000 / rateLimit

    local next = ts + interval;
    local last = redis.call('ZRANGE', retryKey, -1, -1, 'WITHSCORES')

    if type(last) == 'table' and #last > 0 then
        next = tonumber(last[#last]) + interval
    end

    {
        redis.call('ZADD', retryKey, next, next);
        return next - ts;
    }
end
```

add the new member and return
the difference between the current
time and the end of the line



Let's test it

RATE LIMITS

node.js script:

- runs the lua script every 100ms
- if rate limited, schedule a retry based on the returned delay
- periodically check # of requests accepted per second

```
rate limited, retry after: 6077
rate limited, retry after: 6196
rate limited, retry after: 6376
rate limited, retry after: 6472
requests accepted in the last last 1000ms: 5
rate limited, retry after: 6595
rate limited, retry after: 6770
rate limited, retry after: 6869
accepted, waited: 3395
rate limited, retry after: 6968
accepted, waited: 0
accepted, waited: 3536
rate limited, retry after: 6965
accepted, waited: 0
rate limited, retry after: 6994
accepted, waited: 0
rate limited, retry after: 7053
rate limited, retry after: 7196
rate limited, retry after: 7353
rate limited, retry after: 7453
^C
```

Caveats?

```
local rateLimit = 5

local rateKey = KEYS[1]
local retryKey = KEYS[2]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)
redis.call('ZREMRANGEBYSCORE', retryKey, '-inf', ts)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    local interval = 1000 / rateLimit

    local next = ts + interval;
    local last = redis.call('ZRANGE', retryKey, -1, -1, 'WITHSCORES')

    if type(last) == 'table' and #last > 0 then
        next = tonumber(last[#last]) + interval
    end

    redis.call('ZADD', retryKey, next, next);

    return next - ts;
end
```

Caveats?

no limit

cap the retry sorted set

```
local rateLimit = 5

local rateKey = KEYS[1]
local retryKey = KEYS[2]
local ts = tonumber(ARGV[1])

redis.call('ZREMRANGEBYSCORE', rateKey, '-inf', ts - 1000)
redis.call('ZREMRANGEBYSCORE', retryKey, '-inf', ts)

if redis.call('ZCARD', rateKey) < rateLimit then
    redis.call('ZADD', rateKey, ts, ts)
    return 0;
else
    local interval = 1000 / rateLimit

    local next = ts + interval;
    local last = redis.call('ZRANGE', retryKey, -1, -1, 'WITHSCORES')

    if type(last) == 'table' and #last > 0 then
        next = tonumber(last[#last]) + interval
    end

    redis.call('ZADD', retryKey, next, next);

    return next - ts;
end
```

Redis: not just a cache



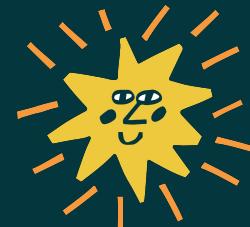
Thank you

 alavers

 alavers

 alavers@zendesk.com

THE BONUS MEAT



bonus

Redis wire protocol

RESP (REdis Serialization Protocol)

TCP 6379

encoded parts strings, integers, arrays, errors
parts terminated with CRLF /r/n

part prefixes:

prefixes:

+ = simple string
- = error
: = integer
\$n = string of n bytes
*n = array of n items

examples:

+OK\r\n
-WRONGTYPE...\r\n
:1\r\n
\$3\r\nget\r\n
*2\r\n:n:1\r\n:n:1\r\n

> get foo
bar



*2\r\n\$3\r\nget\r\n\$3\r\nfoo
\$3\r\nbar



Number of words



**Number of
syllables**



Thank you

 alavers

 alavers

 alavers@zendesk.com